BELNET

The network of knowledge

GRID computing with MPI

BEgrid

The Belgian Grid for Research

BELNET - BEgrid

UCL - INGI – 14th December 2010

BELNET
The network of knowledge

Agenda

- A word on BELNET
- GRID in Europe
- BELNET and GRID
- How it works
  - Virtual Organisations
  - Authentication
  - Searching for resources
  - Sending a job
- Message Passing Interface – MPI
  - MPI principles and API
  - GRID and MPI
- References

The network of knowledge

BELNET

- **BELNET is the Belgian National Research and Education Network (NREN)**
  - **Provide network connectivity to education world**
    - **Universities**
    - **Research centers**
    - **High schools**
    - **…**
  - **Offer services**
    - **« base services »**
    - **GRID**
    - **…**
  - **CERTs – Computer Emergency Response Team**
    - **BELNET CERT**
    - **National CERT (CERT.be)**

BELNET
The network of knowledge

GRID in Europe

- **GRID pushed by EU commission and CERN for LHC needs**
  - ***205 173 cores***
  - **PB of storage**
  - ***33* countries**
  - ***73*M€ in 4 years**
  - ***13 000* researchers**
- **In Belgium mainly used by IIHE and UCL for CMS (High-energy physic)**
- **Resources and authentication distributed worldwide**

BELNET
The network of knowledge

« Central » GRID services

- Belgian « virtual organisations » (*VO*) permission management (*VOMS*)
- Work Management Systems (*WMS*)
- Information Services (*BDII*)
- Monitoring
- GRID Security
  - In collaboration with CERTs
- Support and training
- Bring resources to the GRID ;-)

The network of knowledge

BEgrid

- BEgrid is a collaboration between belgian universities to participate in the GRID
- BELNET acts as coordinator
- Round *1000* cores available
  - + *2000* cores given by the Netherlands (SARA)

BELNET
The network of knowledge

Connection to BELNET

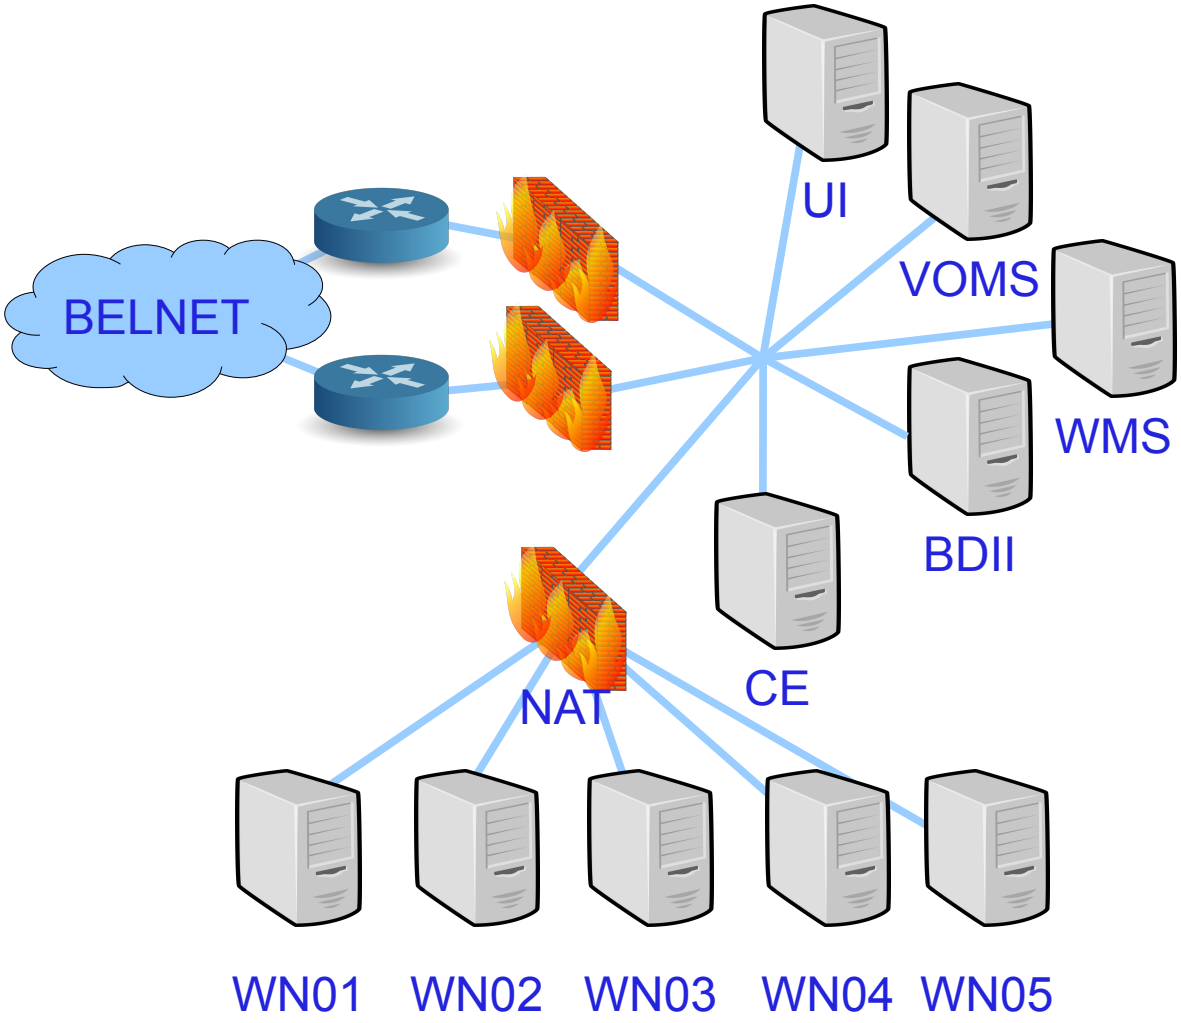

.begrid.be
193.190.113.128/26

BELNET
The network of knowledge

Virtual Organisations

- GRID is not dedicated to a single research field / experiment
- Principe of Virtual Organisations (VO)
- For each VO
  - Each site decide to allow to use resources
  - Priority at site level
  - Access control to softwares, datas...
- Membership to some VO is controlled on VOMS servers (Virtual Organisation Management System)
  - For BELNET: 4 local VOs
    - betest
    - beapps
    - becms and becms-t2

BELNET
The network of knowledge

Security principles

- **Authentication:**
  - Each server has a certificate valid for 2 years
  - Each user has a certificate valid for 1 year
  - Each certificate is signed by a recognized registration authority (RA)
    - BELNET is a recognized RA
- **Authorisation:**
  - Each job is send with a certificate signed by a VOMS server
    - Valid for maximum 24 hours
    - Per default, 8 hours on BELNET Virtual Organisations

BELNET
The network of knowledge

Security principles (2)

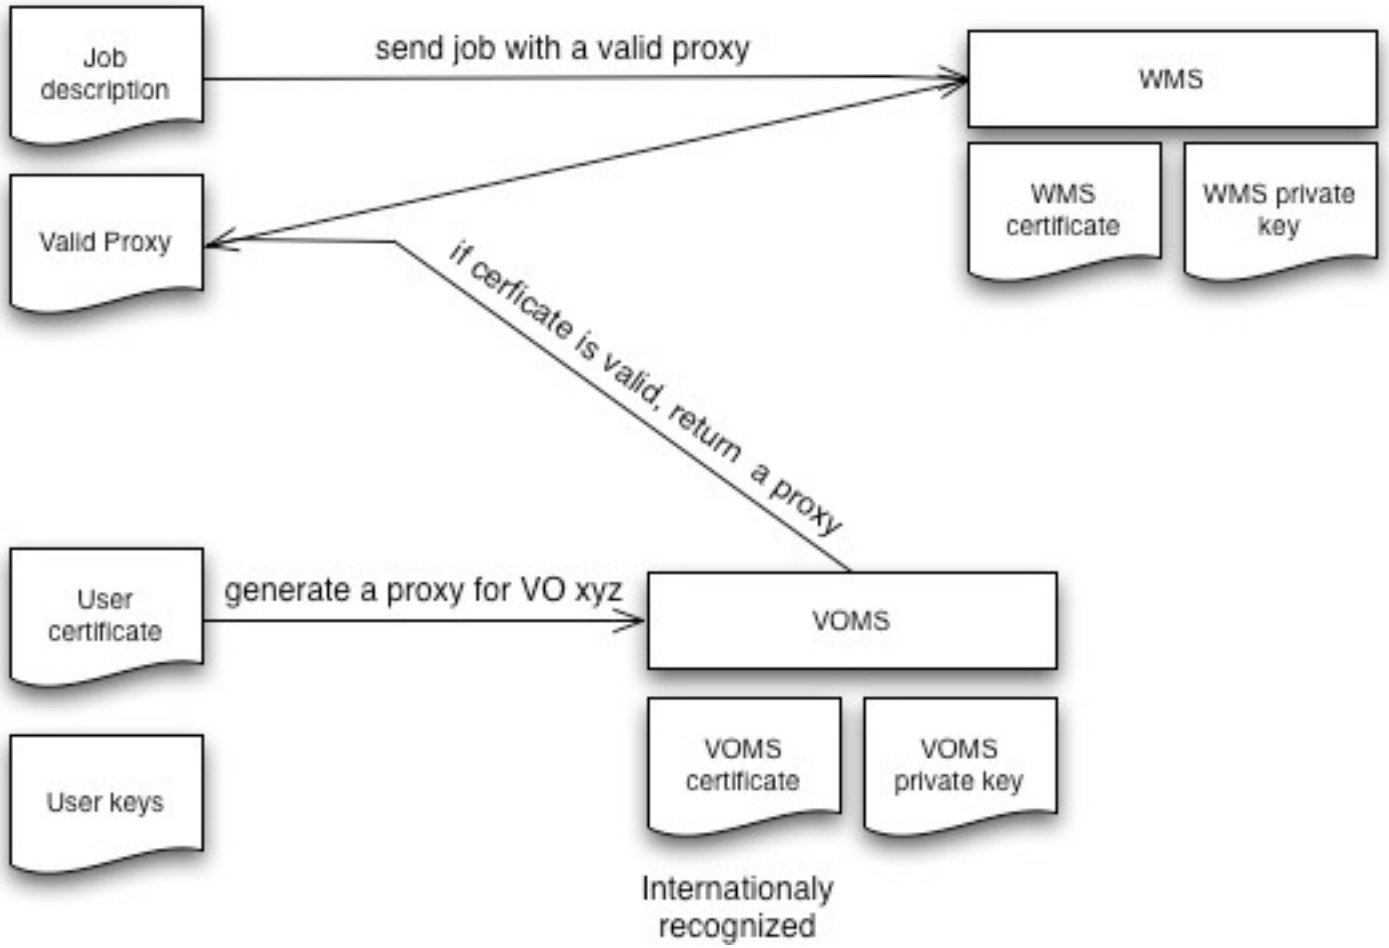

Each server is able to check identity of his pair

BELNET
The network of knowledge

Searching for resources

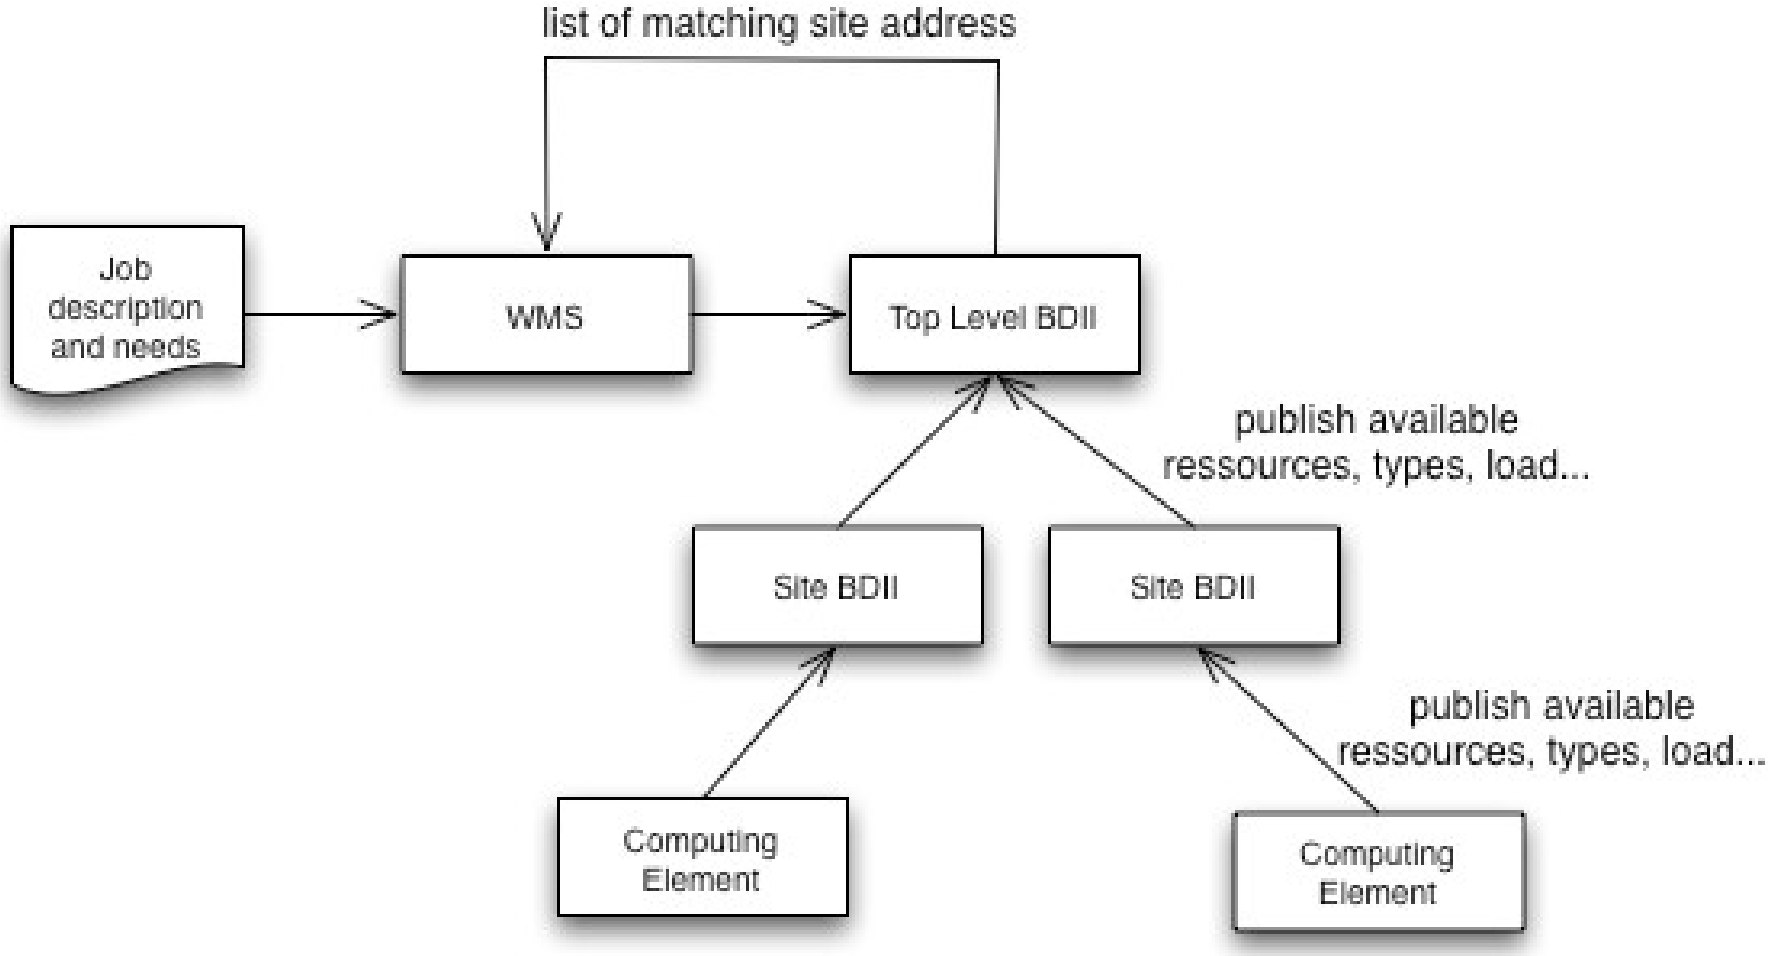

BELNET
The network of knowledge

Global overview

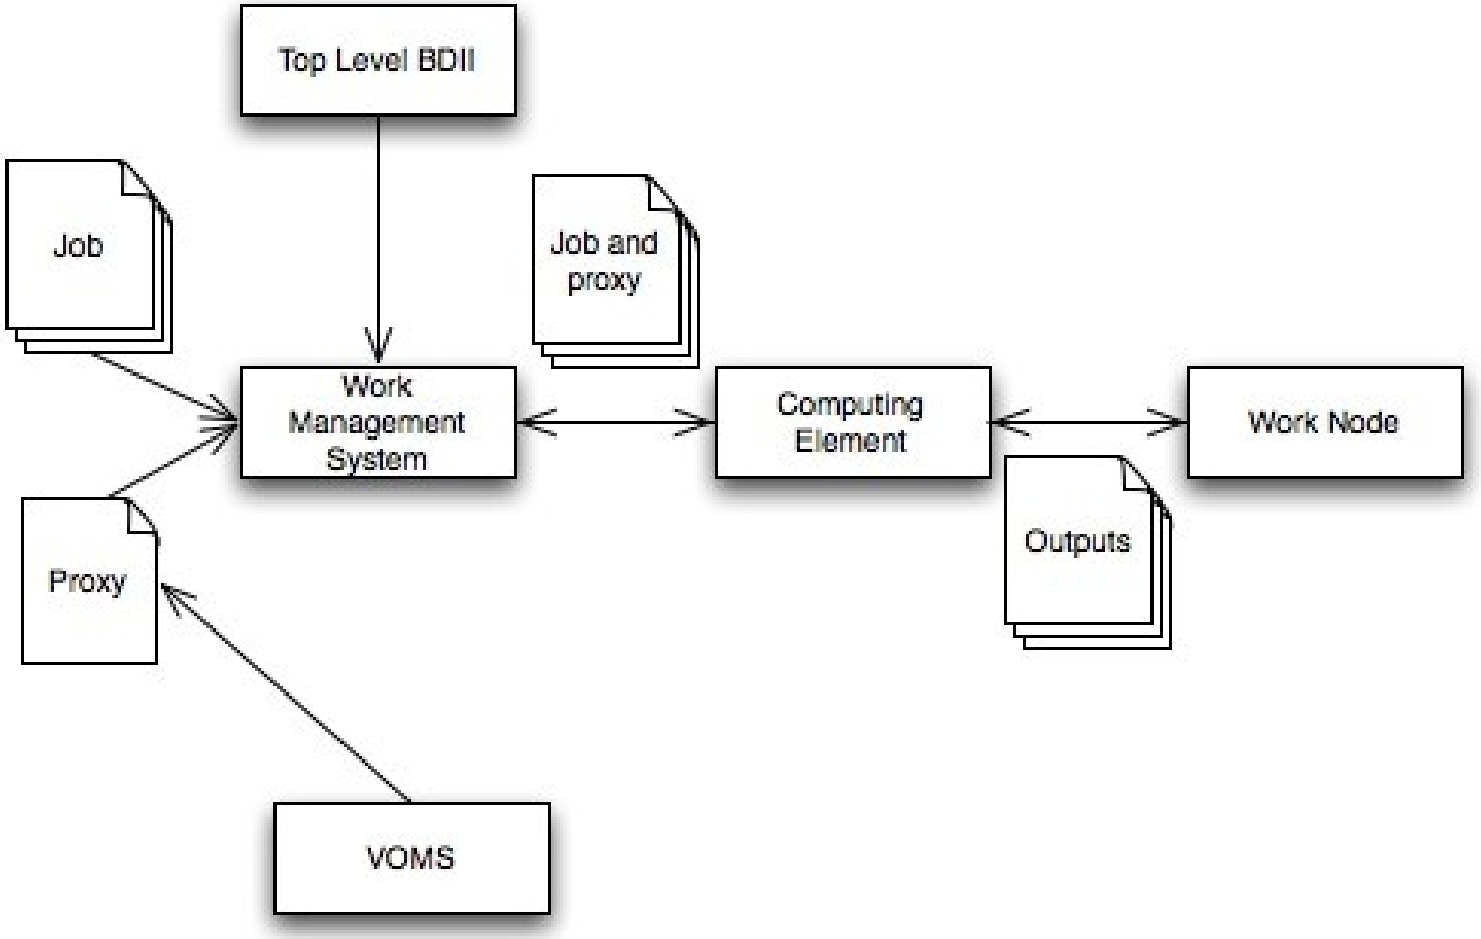

BELNET
The network of knowledge

A simple job

- **Job described using the Job Description Language (JDL)**
- **« Minimum » set of parameters:**
  - **Command / binary to execute**
  - **Input parameters**
  - **Output filename**
  - **Error filename**
  - **Sandbox(es)**

```
Executable = "/bin/echo";
Arguments = "Hello World";
Stdoutput = "message.txt";
StdError = "error.txt";
OutputSandbox = {"message.txt","error.txt"};
```

BELNET
The network of knowledge

Software stack

- Operating Systems – RedHat based
  - ScientificLinux 5.x – Computing resources
  - CentOS 5.x
- Middleware
  - gLite 3.1 and gLite 3.2 (www.glite.org)
- Central management of OS and softwares
  - Quattor (www.quattor.org)
- Regristration Authority
  - OpenTrust (www.opentrust.com)
- Firewalls
  - pfSense (www.pfsense.org)

BELNET
The network of knowledge

Requesting access

1. Request a user certificate
   - At BELNET (https://gridra.belnet.be)
   - At TERENA
   - Any recognized institution
2. Request VO membership
   - Depends on the VO
   - https://voms.begrid.be:8443/vomses/
     - For betest, beapps, becms and becms-t2
3. Request an account on a user interface
   - At BELNET, send a mail to begrid@belnet.be with your public SSH key for ui.begrid.be

BELNET
The network of knowledge

« Defaults » applications

- Compilers / Interpreters
  - gcc – C, C++, Fortran...
  - Perl
  - Python
  - MPI compiler
- Computing environments
  - Octave
  - R
- VO specifics applications
- Other applications
  - Based on needs
  - Based on licenses
  - If needed, request!

BELNET

The network of knowledge

Job control environnent

- Command line interface
  - Typically a SSH on a machine with tools installed
  - ui.begrid.be for instance – not a computing node!
- Web interfaces

BELNET
The network of knowledge

CLI tools – execute a job

- Create a valid proxy
  - voms-proxy-init --voms betest
- Send a job using the proxy
  - glite-wms-job-submit -a myjob.jdl
    - Output an address on the WMS
- Get job status
  - glite-wms-job-status wms_job_address
- Retrieve result
  - glite-wms-job-output -a --dir /path/you/want wms_job_address
- Get matching ressources
  - glite-wms-job-list-match -a myjob.jdl

BELNET
The network of knowledge

Message Passing Interface - MPI

- Language-independent
- Message passing programming
- MPI Goals
  - High performance
  - Scalability
  - Portability
- Two major versions used
  - MPI-1.2
  - MPI-2
- Often one process per processor / core

BELNET
The network of knowledge

MPI-1 and MPI-2

- MPI-1 (1994)
  - Point-to-point communication
  - Collective operations
  - Process groups and topologies
  - Communication contexts
  - Datatype management
- MPI-2 (1997)
  - Dynamic process management
  - File I/O
  - One sided communications
  - Extension of collective operations

BELNET
The network of knowledge

MPI concepts: Communicator

- Connect groups of processes in the MPI session
- In each session
  - Each process receive an independent ID
  - Communicator create a topology
- Communication is
  - Single group intra-communication
  - Bilateral inter-communication
- Could be partitioned

BELNET
The network of knowledge

MPI concepts: Point-to-point

- MPI_SEND
  - Allows one process to send a message to another process
- Blocking or not blocking communication
- « Ready-send »
  - A send could only be made when the matching receive request has been done

BELNET
The network of knowledge

MPI concepts: Collectives

- Communication among all processes in a process group
- MPI_Bcast
  - One process send a message to all other processes in the process group
- MPI_Reduce
  - Take data from all processes
  - Perform some operation
  - Store result
- MPI_Alltoall
  - *n* items exchanged
  - *n*th node receive the *n*th item from each node

Inspired from Wikipedia MPI page (http://www.wikipedia.org)

BELNET
The network of knowledge

MPI concepts: Datatypes

- Pre-defined MPI types for standard types
  - MPI_INT for int
  - MPI_CHAR for char
  - MPI_DOUBLE for double
  - ...

BELNET
The network of knowledge

MPI concepts: Datatypes

- Pre-defined MPI types for standard types
  - MPI_INT for int
  - MPI_CHAR for char
  - MPI_DOUBLE for double
  - ...

BELNET
The network of knowledge

GRID and MPI

- To simplify « GRIDification » of MPI Jobs
  - MPI-Start
    - Portable
    - Permit to enable debug
    - Interface to run MPI job
      - MPI command invisible to user
    - Allow to run MPI job without change to GRID middleware
  - Wrapper
    - Set environment for MPI-Start
    - Call MPI-Start
    - Submitted with the job
  - Hooks
    - Handle compilation
    - Submitted with the job

BELNET

The network of knowledge

Start wrapper

```
#!/bin/bash

MY_EXECUTABLE=`pwd`/$1
MPI_FLAVOR=$2

MPI_FLAVOR_LOWER=`echo $MPI_FLAVOR | tr '[:upper:]' '[:lower:]'`

eval MPI_PATH=`printenv MPI_${MPI_FLAVOR}_PATH`

eval I2G_${MPI_FLAVOR}_PREFIX=$MPI_PATH
export I2G_${MPI_FLAVOR}_PREFIX

touch $MY_EXECUTABLE

export I2G_MPI_APPLICATION=$MY_EXECUTABLE
export I2G_MPI_APPLICATION_ARGS=
export I2G_MPI_TYPE=$MPI_FLAVOR_LOWER
export I2G_MPI_PRE_RUN_HOOK=mpi-hooks.sh
export I2G_MPI_POST_RUN_HOOK=mpi-hooks.sh

export I2G_MPI_START_VERBOSE=1
#export I2G_MPI_START_DEBUG=1

$I2G_MPI_START
```

BELNET
The network of knowledge

Start wrapper

```
export I2G_MPI_APPLICATION
```
- Executable

```
export I2G_MPI_APPLICATION_ARGS
```
- Parameters to give to the executable

```
export I2G_MPI_TYPE
```
- MPI implementation (*OpenMPI* for BEgrid)

```
export I2G_MPI_PRE_RUN_HOOK
export I2G_MPI_POST_RUN_HOOK
```
- Path to hooks

```
export I2G_MPI_START_VERBOSE=1
export I2G_MPI_START_DEBUG=1
```
- Enable verbose / debug modes

BELNET
The network of knowledge

Hooks

```
#!/bin/sh

pre_run_hook () {

  echo "Compiling ${I2G_MPI_APPLICATION}"
  cmd="mpicc ${MPI_MPICC_OPTS} -o ${I2G_MPI_APPLICATION} ${I2G_MPI_APPLICATION}.c"
  echo $cmd
  $cmd
  if [ ! $? -eq 0 ]; then
    echo "Error compiling program.  Exiting..."
    exit 1
  Fi

  echo "Successfully compiled ${I2G_MPI_APPLICATION}"
  return 0
}

post_run_hook () {
  echo "Executing post hook."
  echo "Finished the post hook."

  return 0
}
```

BELNET
The network of knowledge

C job – Hello World

```c
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[]) {

  int numprocs;  /* Number of processors */
  int procnum;   /* Processor number */

  /* Initialize MPI */
  MPI_Init(&argc, &argv);

  /* Find this processor number */
  MPI_Comm_rank(MPI_COMM_WORLD, &procnum);

  /* Find the number of processors */
  MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
  printf ("Hello world! from processor %d out of %d\n", procnum, numprocs);

  /* Shut down MPI */
  MPI_Finalize();
  return 0;
}
```

BELNET
The network of knowledge

Job description

```
JobType = "MPICH";
NodeNumber = 16;
Executable = "mpi-start-wrapper.sh";
Arguments = "mpi-test OPENMPI";
StdOutput = "mpi-test.out";
StdError = "mpi-test.err";
InputSandbox = {"mpi-start-wrapper.sh","mpi-hooks.sh","mpi-test.c"};
OutputSandbox = {"mpi-test.err","mpi-test.out"};
Requirements =
  Member("MPI-START", other.GlueHostApplicationSoftwareRunTimeEnvironment)
  && Member("OPENMPI", other.GlueHostApplicationSoftwareRunTimeEnvironment)
  ;
```

- Ask for a site with MPI-START and OPENMPI
- Job Type set to MPI
- 16 cores claimed

BELNET
The network of knowledge

Ressources

- BELNET trainings
- Hands-on
- BEgrid website
  - http://www.begrid.be
- BEgrid WIKI
  - http://quattorrepository.begrid.be
  - Only available from R&E institutions
- GridCafé
  - http://www.gridcafe.org

BELNET
The network of knowledge

Ressources - MPI

- The Message Passing Interface (MPI) Standard
  - http://www.mcs.anl.gov/research/projects/mpi/
- MPI Documents
  - http://www.mpi-forum.org/docs/docs.html
- Open MPI documentation
  - http://www.open-mpi.org/doc/
- EGEE and EGI projects documentation
  - http://www.eu-egee.org/fileadmin/documents/UseCases/MPIJobs.html
  - https://quattorrepository.begrid.be/trac/centralised-begrid-v6/wiki/MPI_on_the_grid
- Wikipedia

BELNET
The network of knowledge

Questions and answers

Thanks!

?

Feel free to contact us at:
begrid@belnet.be
http://www.begrid.be

BELNET
The network of knowledge

Now... do it yourself :-D

- Compute the value of π using MPI

$$\pi \approx \int_{-1/2}^{1/2} \frac{4}{1+x^2}\, dx$$

Hint:

- Use MPI_Broadcast and MPI_Reduce calls
- Compute by summing rectangles
- Help: http://www.open-mpi.org/doc/v1.5/